



[Click to Take the FREE Imbalanced Classification Crash-Course](#)



How to Fix k-Fold Cross-Validation for Imbalanced Classification

by **Jason Brownlee** on [January 13, 2020](#) in [Imbalanced Classification](#)

Tweet

Share

Share

Last Updated on May 29, 2020

Model evaluation involves using the available dataset to fit a model and estimate its performance when making predictions on unseen examples.

It is a challenging problem as both the training dataset used to fit the model and the test set used to evaluate it must be sufficiently large and representative of the underlying problem so that the resulting estimate of model performance is not too optimistic or pessimistic.

The two most common approaches used for model evaluation are the train/test split and the k-fold cross-validation procedure. Both approaches can be very effective in general, although they can result in misleading results and potentially fail when used on classification problems with a severe class imbalance.

In this tutorial, you will discover how to evaluate classifier models on imbalanced datasets.

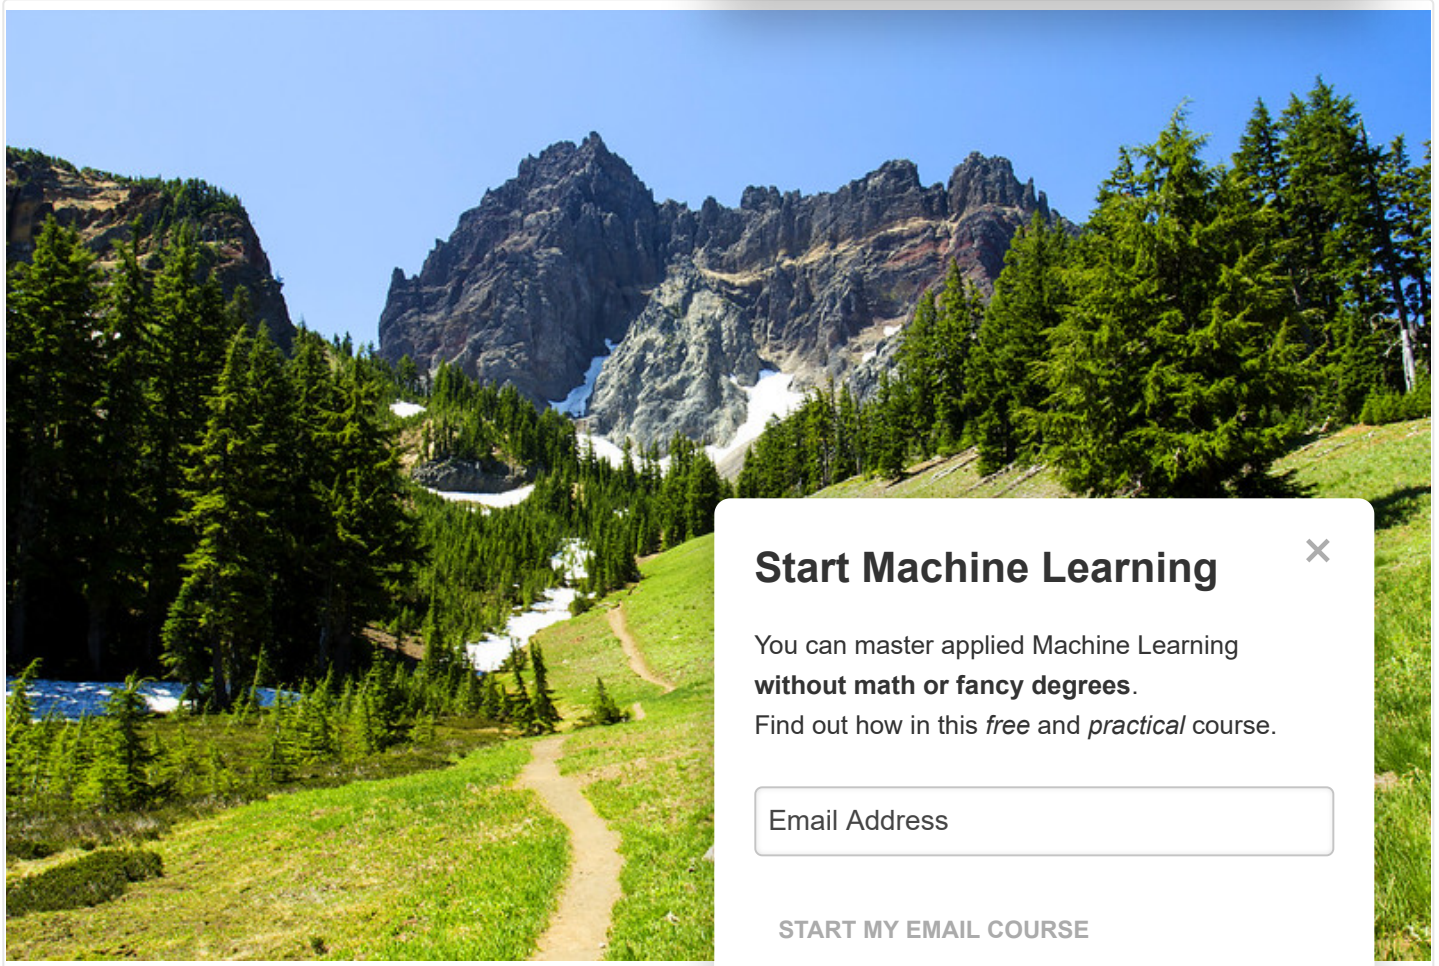
After completing this tutorial, you will know:

- The challenge of evaluating classifiers on datasets using train/test splits and cross-validation.
- How a naive application of k-fold cross-validation and train-test splits will fail when evaluating classifiers on imbalanced datasets.
- How modified k-fold cross-validation and train-test splits can be used to preserve the class distribution in the dataset.

Discover SMOTE, one-class classification, cost-sensitive learning, threshold moving, and much more [in my new book](#), with 30 step-by-step tutorials and full Python source code.

Let's get started.

[Start Machine Learning](#)



Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

How to Use k-Fold Cross-Validation for Imbalanced Classification
Photo by [Bonnie Moreland](#), some rights reserved.

Tutorial Overview

This tutorial is divided into three parts; they are:

1. Challenge of Evaluating Classifiers
2. Failure of k-Fold Cross-Validation
3. Fix Cross-Validation for Imbalanced Classification

Challenge of Evaluating Classifiers

Evaluating a classification model is challenging because we won't know how good a model is until it is used.

Instead, we must estimate the performance of a model using available data where we already have the target or outcome.

Model evaluation involves more than just evaluating a model; it includes testing different data preparation schemes, different learning algorithms, and different hyperparameters for well-performing learning algorithms.

Start Machine Learning

- Model = Data Preparation + Learning Algorithm + Hyperparameters

Ideally, the model construction procedure (data preparation, learning algorithm, and hyperparameters) with the best score (with your chosen metric) can be selected and used.

The simplest model evaluation procedure is to split a dataset into two parts and use one part for training a model and the second part for testing the model. As such, the parts of the dataset are named for their function, train set and test set respectively.

This is effective if your collected dataset is very large and representative of the problem. The number of examples required will differ from problem to problem, but may be thousands, hundreds of thousands, or millions of examples to be sufficient.

A split of 50/50 for train and test would be ideal, although 70/30 or 80/20 for train and test sets.

We rarely have enough data to get an unbiased estimate of a model. Instead, we often have a much smaller dataset and strategies must be used on this dataset.

The most used model evaluation scheme for classifiers is **k-fold cross-validation**.

The **k-fold cross-validation procedure** involves splitting the dataset into k folds. $k-1$ folds are used to train a model, and the holdout k th fold is used to evaluate the model. Each of the folds is given an opportunity to be used as the holdout test set. A total of k models are fit and evaluated, and the performance of the model is calculated as the mean of these runs.

The procedure has been shown to give a less optimistic estimate of model performance on small training datasets than a single train/test split. A value of $k=10$ has been shown to be effective across a wide range of dataset sizes and model types.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Want to Get Started With Imbalance Classification?

Take my free 7-day email crash course now (with sample code).

Click to sign-up and also get a free PDF Ebook version of the course.

Download Your **FREE** Mini-Course

Failure of k-Fold Cross-Validation

Sadly, the k-fold cross-validation is not appropriate for

Start Machine Learning



A 10-fold cross-validation, in particular, the most commonly used error-estimation method in machine learning, can easily break down in the case of class imbalances, even if the skew is less extreme than the one previously considered.

— Page 188, [Imbalanced Learning: Foundations, Algorithms, and Applications](#), 2013.

The reason is that the data is split into k -folds with a uniform probability distribution.

This might work fine for data with a balanced class distribution, but when the distribution is severely skewed, it is likely that one or more folds will have few or no examples from the minority class. This means that some or perhaps many of the model evaluations will be misleading, as the model need only predict the majority class correctly.

We can make this concrete with an example.

First, we can define a dataset with a 1:100 minority to

This can be achieved using the `make_classification()` function, specifying the number of examples (1,000), the number of classes

```
1 # generate 2 class dataset
2 X, y = make_classification(n_samples=1000, n_c
```

The example below generates the synthetic binary class distribution.

```
1 # create a binary classification dataset
2 from numpy import unique
3 from sklearn.datasets import make_classification
4 # generate 2 class dataset
5 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], flip_y=0, random_
6 # summarize dataset
7 classes = unique(y)
8 total = len(y)
9 for c in classes:
10     n_examples = len(y[y==c])
11     percent = n_examples / total * 100
12     print('> Class=%d : %d/%d (%.1f%%)' % (c, n_examples, total, percent))
```

Running the example creates the dataset and summarizes the number of examples in each class.

By setting the `random_state` argument, it ensures that we get the same randomly generated examples each time the code is run.

```
1 > Class=0 : 990/1000 (99.0%)
2 > Class=1 : 10/1000 (1.0%)
```

A total of 10 examples in the minority class is not many. If we used 10-folds, we would get one example in each fold in the ideal case, which is not enough to train a model. For demonstration purposes, we will use 5-folds.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning

In the ideal case, we would have 10/5 or two examples in each fold, meaning 4*2 (8) folds worth of examples in a training dataset and 1*2 folds (2) in a given test dataset.

First, we will use the `KFold` class to randomly split the dataset into 5-folds and check the composition of each train and test set. The complete example is listed below.

```
1 # example of k-fold cross-validation with an imbalanced dataset
2 from sklearn.datasets import make_classification
3 from sklearn.model_selection import KFold
4 # generate 2 class dataset
5 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], flip_y=0, random_
6 kfold = KFold(n_splits=5, shuffle=True, random_state=1)
7 # enumerate the splits and summarize the distributions
8 for train_ix, test_ix in kfold.split(X):
9     # select rows
10    train_X, test_X = X[train_ix], X[test_ix]
11    train_y, test_y = y[train_ix], y[test_ix]
12    # summarize train and test composition
13    train_0, train_1 = len(train_y[train_y==0]), len(train_y[train_y==1])
14    test_0, test_1 = len(test_y[test_y==0]), len(test_y[test_y==1])
15    print('>Train: 0=%d, 1=%d, Test: 0=%d, 1=%d'
```

Running the example creates the same dataset and enumerates the splits, showing the distribution for both the train and test sets.

We can see that in this case, there are some splits that are good, such as 9/1 (optimist

Evaluating a model on these splits of the data would not give a reliable estimate of performance.

```
1 >Train: 0=791, 1=9, Test: 0=199, 1=1
2 >Train: 0=793, 1=7, Test: 0=197, 1=3
3 >Train: 0=794, 1=6, Test: 0=196, 1=4
4 >Train: 0=790, 1=10, Test: 0=200, 1=0
5 >Train: 0=792, 1=8, Test: 0=198, 1=2
```

We can demonstrate a similar issue exists if we use a simple train/test split of the dataset, although the issue is less severe.

We can use the `train_test_split()` function to create a 50/50 split of the dataset and, on average, we would expect five examples from the minority class to appear in each dataset if we performed this split many times.

The complete example is listed below.

```
1 # example of train/test split with an imbalanced dataset
2 from sklearn.datasets import make_classification
3 from sklearn.model_selection import train_test_split
4 # generate 2 class dataset
5 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], flip_y=0, random_
6 # split into train/test sets with same class ratio
7 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2)
8 # summarize
9 train_0, train_1 = len(trainy[trainy==0]), len(trainy[trainy==1])
10 test_0, test_1 = len(testy[testy==0]), len(testy[testy==1])
11 print('>Train: 0=%d, 1=%d, Test: 0=%d, 1=%d'
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Running the example creates the same dataset as before and splits it into a random train and test split.

In this case, we can see only three examples of the minority class are present in the training set, with seven in the test set.

Evaluating models on this split would not give them enough examples to learn from, too many to be evaluated on, and likely give poor performance. You can imagine how the situation could be worse with an even more severe random split.

```
1 >Train: 0=497, 1=3, Test: 0=493, 1=7
```

Fix Cross-Validation for Imbalanced Classification

The solution is to not split the data randomly when using cross-validation.

Specifically, we can split a dataset randomly, although the distribution in each subset. This is called stratification. Stratification, by class, is used to control the sampling process.

For example, we can use a version of k-fold cross-validation that ensures the distribution in each fold. It is called stratified k-fold cross-validation. In each split of the data to match the distribution in the original dataset.

“... it is common, in the case of class imbalances in particular, to use stratified 10-fold cross-validation, which ensures that the proportion of positive to negative examples found in the original distribution is respected in all the folds.

— Page 205, [Imbalanced Learning: Foundations, Algorithms, and Applications](#), 2013.

We can make this concrete with an example.

We can stratify the splits using the [StratifiedKFold](#) class that supports stratified k-fold cross-validation as its name suggests.

Below is the same dataset and the same example with the stratified version of cross-validation.

```
1 # example of stratified k-fold cross-validation with an imbalanced dataset
2 from sklearn.datasets import make_classification
3 from sklearn.model_selection import StratifiedKFold
4 # generate 2 class dataset
5 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], flip_y=0, random_state=1)
6 kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)
7 # enumerate the splits and summarize the distributions
8 for train_ix, test_ix in kfold.split(X, y):
9     # select rows
10    train_X, test_X = X[train_ix], X[test_ix]
11    train_y, test_y = y[train_ix], y[test_ix]
12    # summarize train and test composition
13    train_0, train_1 = len(train_y[train_y==0]), len(train_y[train_y==1])
14    test_0, test_1 = len(test_y[test_y==0]), len(test_y[test_y==1])
```

Start Machine Learning

```
15 print('>Train: 0=%d, 1=%d, Test: 0=%d, 1=%d' % (train_0, train_1, test_0, test_1))
```

Running the example generates the dataset as before and summarizes the class distribution for the train and test sets for each split.

In this case, we can see that each split matches what we expected in the ideal case.

Each of the examples in the minority class is given one opportunity to be used in a test set, and each train and test set for each split of the data has the same class distribution.

```
1 >Train: 0=792, 1=8, Test: 0=198, 1=2
2 >Train: 0=792, 1=8, Test: 0=198, 1=2
3 >Train: 0=792, 1=8, Test: 0=198, 1=2
4 >Train: 0=792, 1=8, Test: 0=198, 1=2
5 >Train: 0=792, 1=8, Test: 0=198, 1=2
```

This example highlights the need to first select a value where there are a sufficient number of examples in the train and test sets. The number of examples from the minority class in the test set is probably too few.

It also highlights the requirement to use stratified *k*-fold cross-validation to preserve the class distribution in the train and test sets.

We can also use a stratified version of a train/test split.

This can be achieved by setting the “*stratify*” argument to the “*y*” variable containing the target variable from the dataset. From this, the function will determine the desired class distribution and ensure that the train and test sets both have this distribution.

We can demonstrate this with a worked example, listed below.

```
1 # example of stratified train/test split with an imbalanced dataset
2 from sklearn.datasets import make_classification
3 from sklearn.model_selection import train_test_split
4 # generate 2 class dataset
5 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], flip_y=0, random_state=1)
6 # split into train/test sets with same class ratio
7 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
8 # summarize
9 train_0, train_1 = len(trainy[trainy==0]), len(trainy[trainy==1])
10 test_0, test_1 = len(testy[testy==0]), len(testy[testy==1])
11 print('>Train: 0=%d, 1=%d, Test: 0=%d, 1=%d' % (train_0, train_1, test_0, test_1))
```

Running the example creates a random split of the dataset into training and test sets, ensuring that the class distribution is preserved, in this case leaving five examples in each dataset.

```
1 >Train: 0=495, 1=5, Test: 0=495, 1=5
```

Further Reading

This section provides more resources on the topic if you are looking to go deeper.

Start Machine Learning

Tutorials

- [A Gentle Introduction to k-fold Cross-Validation](#)

Books

- [Imbalanced Learning: Foundations, Algorithms, and Applications](#), 2013.

API

- `sklearn.model_selection.KFold` API.
- `sklearn.model_selection.StratifiedKFold` API.
- `sklearn.model_selection.train_test_split` API.

Summary

In this tutorial, you discovered how to evaluate classifiers

Specifically, you learned:

- The challenge of evaluating classifiers on dataset
- How a naive application of k-fold cross-validation classifiers on imbalanced datasets.
- How modified k-fold cross-validation and train-test in the dataset.

Do you have any questions?


Ask your questions in the comments below and I will do my best to answer.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Get a Handle on Imbalanced Classification!



Develop Imbalanced Learning Models in Minutes

...with just a few lines of python code

Discover how in my new Ebook:
[Imbalanced Classification with Python](#)

It provides **self-study tutorials** and **end-to-end projects** on:
Performance Metrics, Undersampling Methods, SMOTE, Threshold Moving, Probability Calibration, Cost-Sensitive Algorithms
and much more...

Bring Imbalanced Classification Methods to Your Machine Learning Projects

Start Machine Learning

SEE WHAT'S INSIDE

Tweet

Share

Share

**About Jason Brownlee**

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

Start Machine Learning ×

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

< A Gentle Introduction to Probability Metrics for Imbalanced Classification
What Is t

26 Responses to *How to Fix k-Fold Cross-Validation for Imbalanced Classification*



Sadeqh Arefnezhad January 14, 2020 at 1:05 am #

REPLY ↩

Thank you!



Jason Brownlee January 14, 2020 at 7:24 am #

REPLY ↩

You're welcome!



Markus January 15, 2020 at 10:53 pm #

REPLY ↩

As always, KUDOS!



Jason Brownlee January 16, 2020 at 6:16 am #

REPLY ↩

Thanks!

Start Machine Learning



avinash January 25, 2020 at 1:13 am #

REPLY ↩

$\text{train_X, test_y} = X[\text{train_ix}], X[\text{test_ix}]$ should be $\text{train_X, test_X} = X[\text{train_ix}], X[\text{test_ix}]$



Jason Brownlee January 25, 2020 at 8:38 am #

REPLY ↩

Thanks, fixed!



Sachin February 15, 2020 at 1:31 am #

Hello, Great article. Thanks.

My problem is that in my original dataset, I have 1% responders and 50% non responders. After this I have validation on sampled data (not real data)? Also, I am test dataset.

Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



Jason Brownlee February 15, 2020 at 6:34 am #

No.

The resampling must happen within the cross-validation folds. See examples here: <https://machinelearningmastery.com/smote-oversampling-for-imbalanced-classification/>



Sachin February 16, 2020 at 2:57 am #

REPLY ↩

Thanks. Even in that example the SMOTE and undersampling of majority class is done before stratified k-fold validation is applied on X and y. This means the validation is done on balanced data. Please correct if I am wrong. Thanks again.



Jason Brownlee February 16, 2020 at 6:13 am #

REPLY ↩

Not true.

We use a pipeline to ensure that data sampling occurs within the cross-validation procedure.

Perhaps re-read the tutorial?

Start Machine Learning



Sachin February 17, 2020 at 9:08 pm #

REPLY ↩

According to the example, this is the pipeline:

```
steps = [('over', over), ('under', under), ('model', model)]  
pipeline = Pipeline(steps=steps)
```

and then this step does the cross validation:

```
scores = cross_val_score(pipeline, X, y, scoring='roc_auc', cv=cv, n_jobs=-1)
```

doesn't this mean that pipeline is applied to original dataset (X and y) – which means first first oversampling by SMOTE, then random undersampling and then the model is applied and validated with ROC curve?



Jason Brownlee February 18, 2020 at 6:19 am #

Yes, but the over/under sampling is applied



faezeh March 25, 2020 at 2:08 am #

Hello,

What is the disadvantage(s) of the method s
Stratified cross validation ?



Jason Brownlee March 25, 2020 at 6:34 am #

REPLY ↩

It is computationally more expensive, slightly.

It is only for classification, not regression.



faezeh March 25, 2020 at 4:39 pm #

REPLY ↩

Thank you



Jason Brownlee March 26, 2020 at 7:49 am #

REPLY ↩

You're welcome.

Start Machine Learning



You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



Daniele April 22, 2020 at 12:30 am #

REPLY ↩

Hi, great post! In the case of GridSearchCV, is (Stratified)KFolds implicit? This is an example:

```
gs_clf = GridSearchCV(clf_pipe, param_grid=params, verbose=0, cv=5, n_jobs=-1)
```

Thanks for your reply!



Jason Brownlee April 22, 2020 at 5:59 am #

REPLY ↩

Thanks.

I think so. It is better to be explicit with your cv method.



ghizlan April 26, 2020 at 12:24 pm #

Hello Jason, great post thank you !

One question

I have a data of 100 samples and 10 features. I want to split it into two sets (variables)

method 1: I divide my data into 80% for training using k-fold cross validation and then validate the model on the unseen data(20%)

Method 2: I use all my data to fit the model using k-fold cross validation.

Between the two methods, which is the true one?

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



Jason Brownlee April 27, 2020 at 5:25 am #

REPLY ↩

They are both viable/true, use the method that you believe will give you the most robust estimate of model performance for your specific project.



ghizlan April 27, 2020 at 7:17 am #

REPLY ↩

Thank you, Jason



Jason Brownlee April 27, 2020 at 7:34 am #

REPLY ↩

You're welcome.

Start Machine Learning



Grzegorz Kępisty May 7, 2020 at 4:53 pm #

REPLY ↩

It seems that “stratify” flag from `train_test_split()` is very useful. What is more, almost any real classification problem is imbalanced. Then I guess that this strategy shall be always used as a default? (as there is nothing to lose, right?)



Jason Brownlee May 8, 2020 at 6:23 am #

REPLY ↩

Exactly right on all points!

I guess the only practical down side is a slight com



Saily Shah May 21, 2020 at 1:18 am #

Hi Jason. Wonderful article:-). I just had one c
splitting of the imbalanced data in train and test dataset
maintained in training and test dataset. In stratified k-f
“training” dataset is divided into `n_folds`, it maintains th
we are handling imbalanced data, we need to use “stratify” flag in `train_test_split` and also use stratified k
fold cv”. Am i correct?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



Jason Brownlee May 21, 2020 at 6:20 am #

REPLY ↩

Thanks.

Correct.

Leave a Reply

Start Machine Learning

Name (required)

Email (will not be published) (required)

Website

SUBMIT COMMENT


**Welcome!**

My name is *Jason Brownlee* PhD, and I help
[Read more](#)

Start Machine Learning ×

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Never miss a tutorial:**Picked for you:**[8 Tactics to Combat Imbalanced Classes in Your Machine Learning Dataset](#)[SMOTE for Imbalanced Classification with Python](#)[Imbalanced Classification With Python \(7-Day Mini-Course\)](#)[A Gentle Introduction to Threshold-Moving for Imbalanced Classification](#)[Step-By-Step Framework for Imbalanced Classification Projects](#)

Loving the

Start Machine Learning

The [Imbalanced Classification with Python](#) EBook is where I keep the ***Really Good*** stuff.

SEE WHAT'S INSIDE

© 2020 Machine Learning Mastery Pty. Ltd. All Rights Reserved.

Address: PO Box 206, Vermont Victoria 3133, Australia. | ACN: 626 223 336.

[LinkedIn](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#) | [Sitemap](#) | [Search](#)

Start Machine Learning ×

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning